

The Lind Virtual Machine

Userspace system call emulation and auditing

Joey Pabalinas

contracted for implementation by

Professor Justin Cappos

Secure Systems Laboratory, New York University

December 20, 2018

Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy

1 General overview of the project

- Purpose
- Architecture
- System Call Anatomy

Purpose

The Lind
Virtual
Machine

Joey
Pabalinas

Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy

Lind is a fork of Google's NaCl.

- 1 Security goals fall somewhere in-between a virtual machine and a container
- 2 All system calls are trapped and handled by Lind
 - Safe System Calls - Python POSIX Compatibility Layer
 - Low-level System Calls - NaCl Internal Emulation
 - Unsafe System Calls - Return Unimplemented

Architecture

The Lind
Virtual
Machine

Joey
Pabalinas

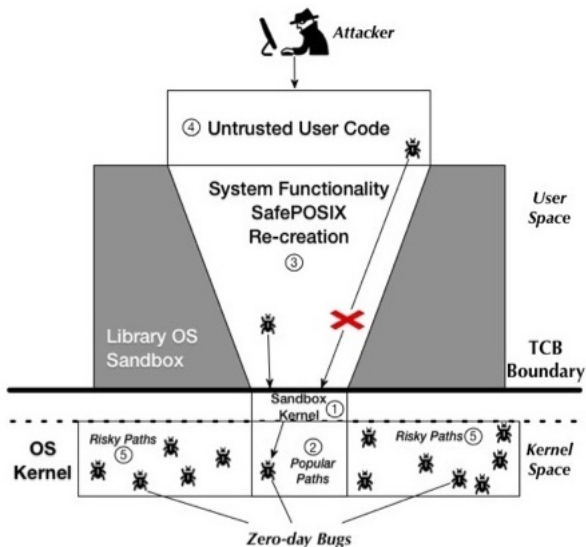
Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy



System Call Anatomy

The Lind
Virtual
Machine

Joey
Pabalinas

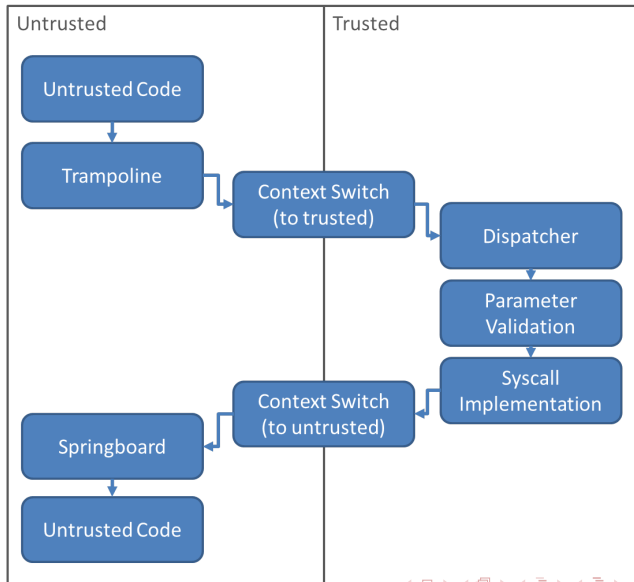
Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy



Debugging Segfaults Inside Lind

```
$ lind -g /bin/bash -c '/hello;_/hello'
```

```
Starting program: sel_ldr [...]  
Thread 7 "sel_ldr" received signal  
SIGSEGV, Segmentation fault.
```

```
0x0000690501297143 in ?? ()
```

Our Incredibly Useful Backtrace

```
@ [#7:None()] backtrace
```

```
#0 0x0000690501297143 in ?? ()
```

```
#1 0x0000000000000000 in ?? ()
```

This ASM Doesn't Seem Right...

```
@ [#7:None()] x/10i $rip
```

```
=> 0x690501297143:  lock cmpxchg %edx,(%r15,%r13,1)
0x690501297149:  test  %eax,%eax
0x69050129714b:  je     0x6905012971a0
0x69050129714d:  nopl   0x0(%rax)
0x690501297151:  data16 nopw %cs:0x0(%rax,%rax,1)
0x690501297160:  mov    %r13d,%edi
0x690501297163:  mov    %r12d,%esi
0x690501297166:  xor    %ebx,%ebx
0x690501297168:  nopl   0x0(%rax)
0x69050129716c:  data16 nopw %cs:0x0(%rax,%rax,1)
```


The Fault Address

```
@ [#7:None()] print $_siginfo._sifields._sigfault.si_addr  
$5 = (void *) 0x6905000368e0
```

Conclusion: The Problem Address: $r15 + r13$

@ [#7:None()] info registers

rax	0x0	0
rbx	0x10057910	268794128
rcx	0x1297120	19493152
[...]		
r13	0x368e0	223456
r14	0x0	0
r15	0x690500000000	115470195752960
rip	0x690501297143	0x690501297143
eflags	0x10246	[PF ZF IF RF]

Conclusion: Memory Not Writable

```
$ grep 6905000 /proc/$(pidof sel_ldr)/maps
```

```
68fb00000000 -690500000000 —p 00000000 00:00 0
690500000000 -690500010000 rw-p 00000000 00:00 0
=> 690500010000 -690510020000 r-xp 00000000 00:00 0
```

Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy

Questions?

Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy

Questions?

O..

Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy

Questions?

O..

.O.

Questions?

O..

.O.

[1] 3286 segmentation fault (core
dumped)

Outline

General
overview of
the project

Purpose

Architecture

System Call Anatomy

Thank you.